

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

```
In [3]: housing.keys()
```

```
Out[3]: dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names', 'DESCR'])
```

```
In [4]: dataset=pd.DataFrame(housing.data,columns=housing.feature_names)
```

```
In [5]: dataset.head()
```

```
Out[5]:   MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  Longitude
0     8.3252      41.0    6.984127    1.023810      322.0    2.555556    37.88   -122.23
1     8.3014      21.0    6.238137    0.971880     2401.0    2.109842    37.86   -122.22
2     7.2574      52.0    8.288136    1.073446      496.0    2.802260    37.85   -122.24
3     5.6431      52.0    5.817352    1.073059      558.0    2.547945    37.85   -122.25
4     3.8462      52.0    6.281853    1.081081      565.0    2.181467    37.85   -122.25
```

```
In [6]: dataset['price']=housing.target
```

```
In [7]: print(dataset['price'])
```

```
0        4.526
1        3.585
2        3.521
3        3.413
4        3.422
...
20635    0.781
20636    0.771
20637    0.923
20638    0.847
20639    0.894
Name: price, Length: 20640, dtype: float64
```

```
In [8]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   MedInc       20640 non-null   float64
 1   HouseAge     20640 non-null   float64
 2   AveRooms     20640 non-null   float64
 3   AveBedrms    20640 non-null   float64
 4   Population   20640 non-null   float64
 5   AveOccup     20640 non-null   float64
 6   Latitude     20640 non-null   float64
 7   Longitude    20640 non-null   float64
 8   price        20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB
```

```
In [9]: dataset.isnull().sum()
```

```
Out[9]: MedInc      0
HouseAge     0
AveRooms     0
AveBedrms    0
Population   0
AveOccup     0
Latitude     0
Longitude    0
price        0
dtype: int64
```

```
In [10]: dataset.describe()
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	L
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	3.070655	35.
std	1.899822	12.585558	2.474173	0.473911	1132.462122	10.386050	2.
min	0.499900	1.000000	0.846154	0.333333	3.000000	0.692308	32.
25%	2.563400	18.000000	4.440716	1.006079	787.000000	2.429741	33.
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	2.818116	34.
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	3.282261	37.
max	15.000100	52.000000	141.909091	34.066667	35682.000000	1243.333333	41.

```
In [11]: dataset.corr()
```

Out[11]:

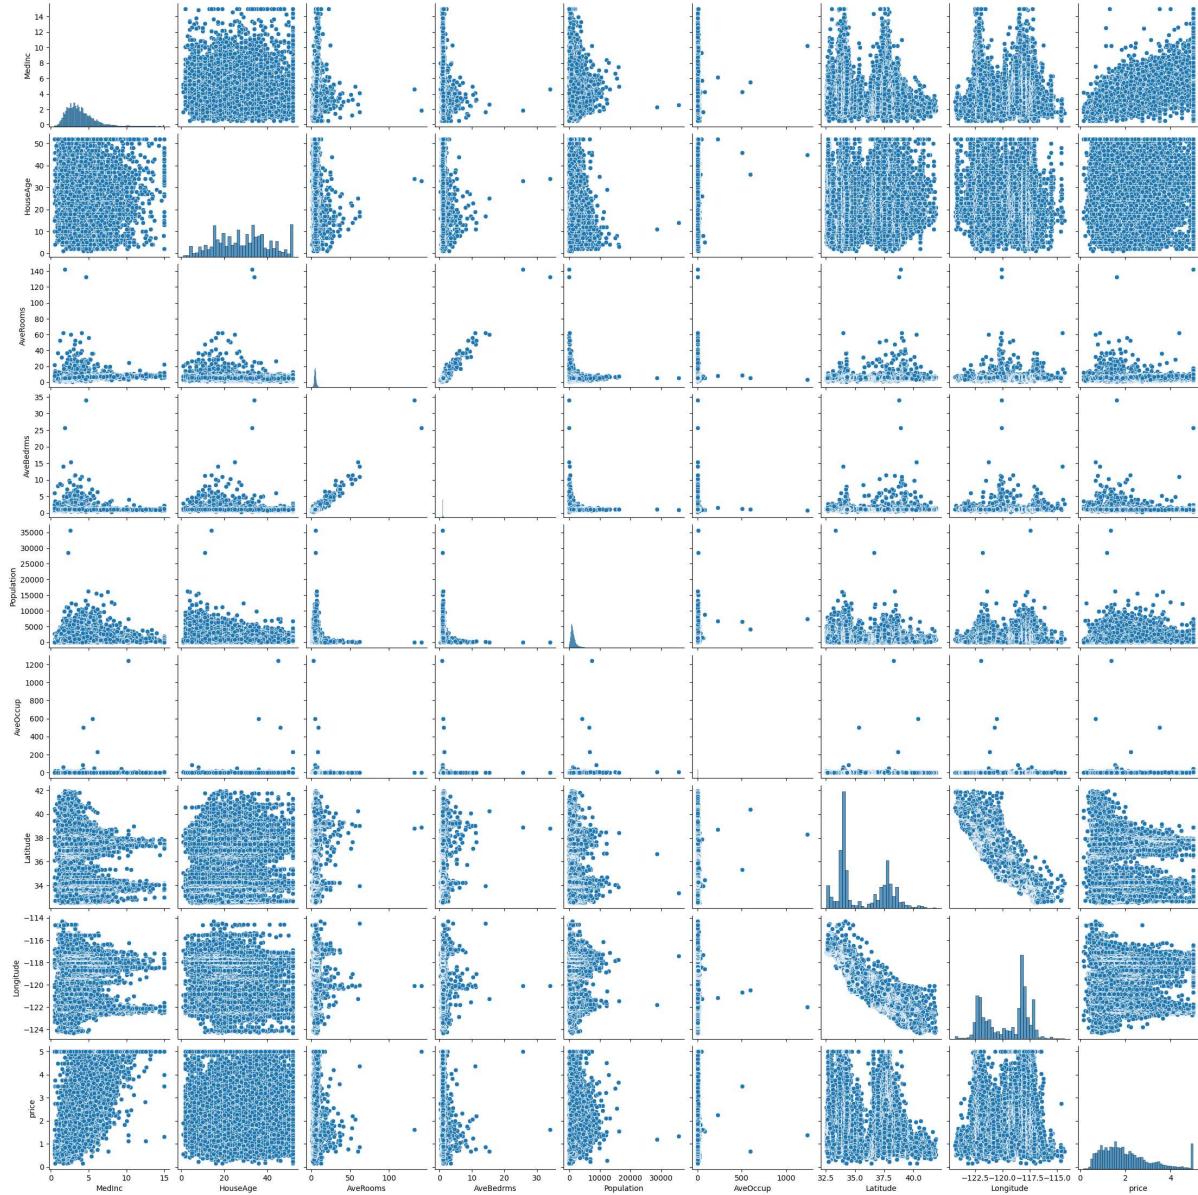
	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Long
MedInc	1.000000	-0.119034	0.326895	-0.062040	0.004834	0.018766	-0.079809	-0.015176
HouseAge	-0.119034	1.000000	-0.153277	-0.077747	-0.296244	0.013191	0.011173	-0.108197
AveRooms	0.326895	-0.153277	1.000000	0.847621	-0.072213	-0.004852	0.106389	-0.027540
AveBedrms	-0.062040	-0.077747	0.847621	1.000000	-0.066197	-0.006181	0.069721	0.013344
Population	0.004834	-0.296244	-0.072213	-0.066197	1.000000	0.069863	-0.108785	0.099773
AveOccup	0.018766	0.013191	-0.004852	-0.006181	0.069863	1.000000	0.002366	0.002476
Latitude	-0.079809	0.011173	0.106389	0.069721	-0.108785	0.002366	1.000000	-0.924664
Longitude	-0.015176	-0.108197	-0.027540	0.013344	0.099773	0.002476	-0.924664	1.000000
price	0.688075	0.105623	0.151948	-0.046701	-0.024650	-0.023737	-0.144160	-0.027540

In [12]: `dataset.columns`

Out[12]: `Index(['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude', 'price'],
 dtype='object')`

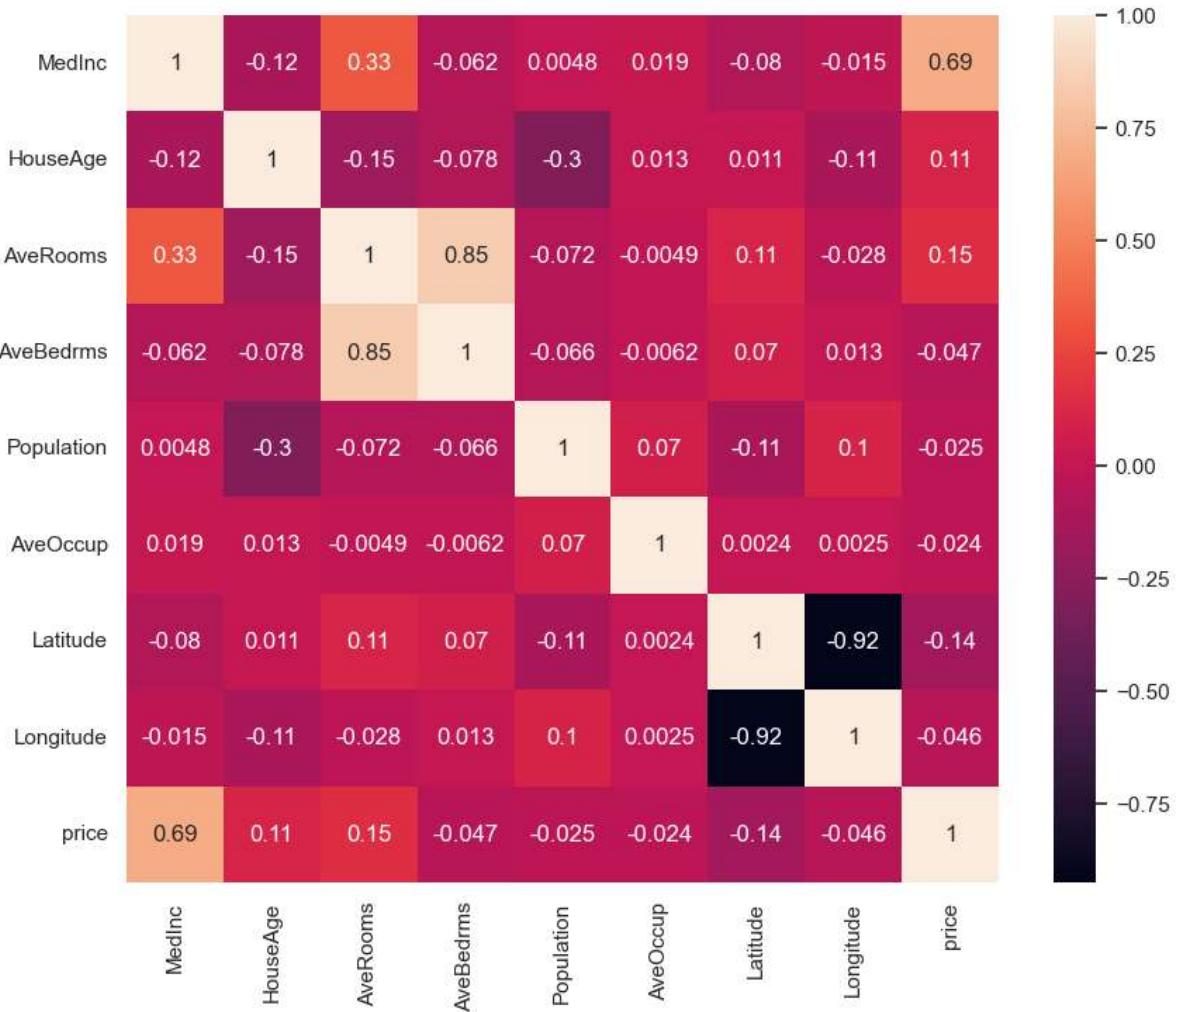
In [13]: `sns.pairplot(dataset)`

Out[13]: `<seaborn.axisgrid.PairGrid at 0x244d14de3b0>`



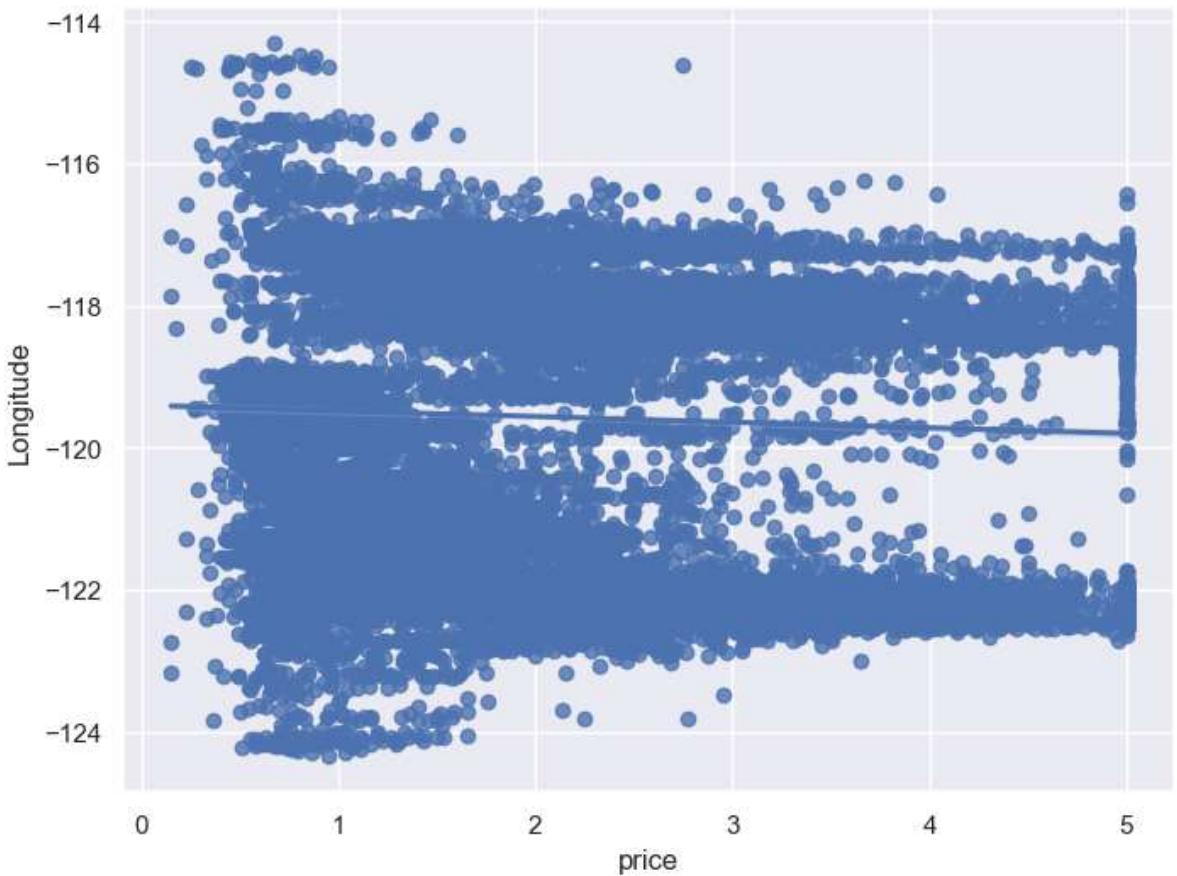
```
In [14]: sns.set(rc={'figure.figsize':(10,8)})  
sns.heatmap(dataset.corr(), annot=True )
```

```
Out[14]: <AxesSubplot: >
```



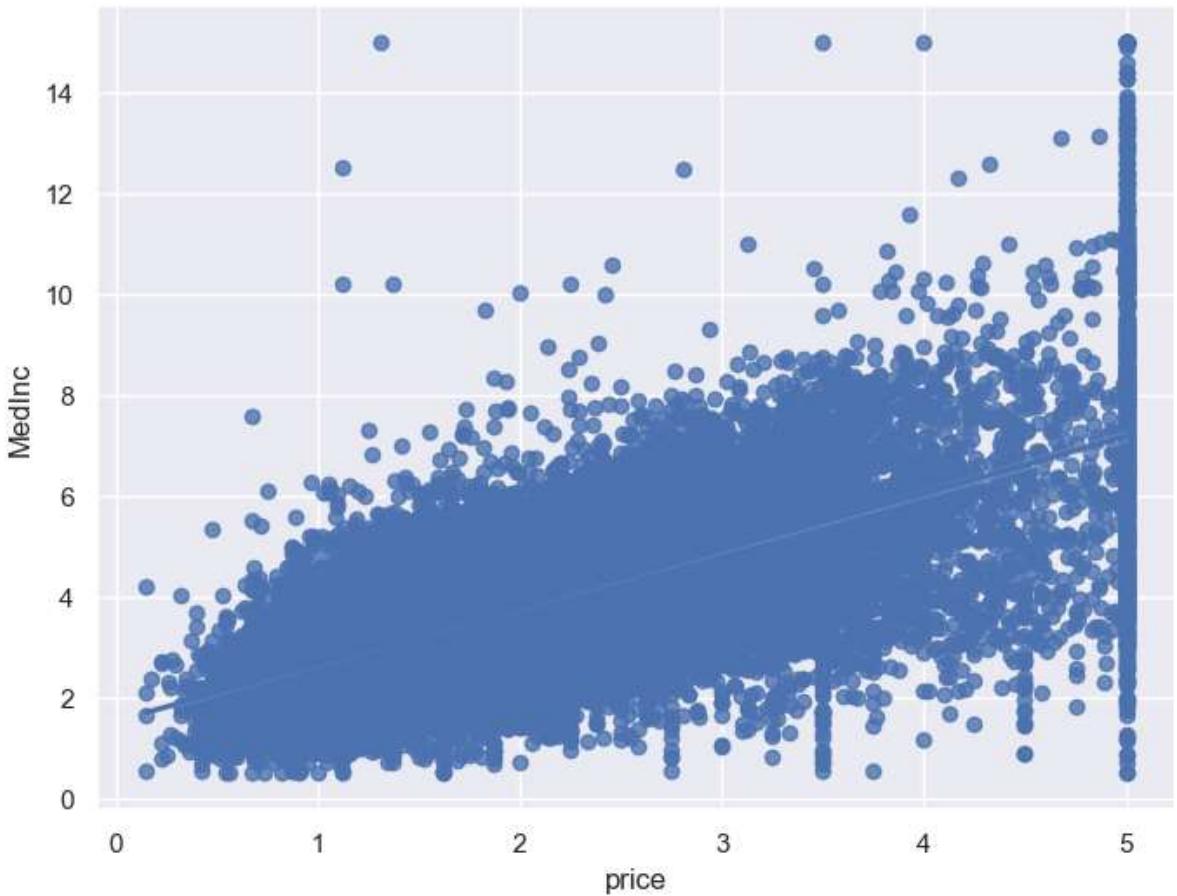
```
In [15]: sns.set(rc={'figure.figsize':(8,6)})
sns.regplot(x='price',y='Longitude',data=dataset )
```

```
Out[15]: <AxesSubplot: xlabel='price', ylabel='Longitude'>
```



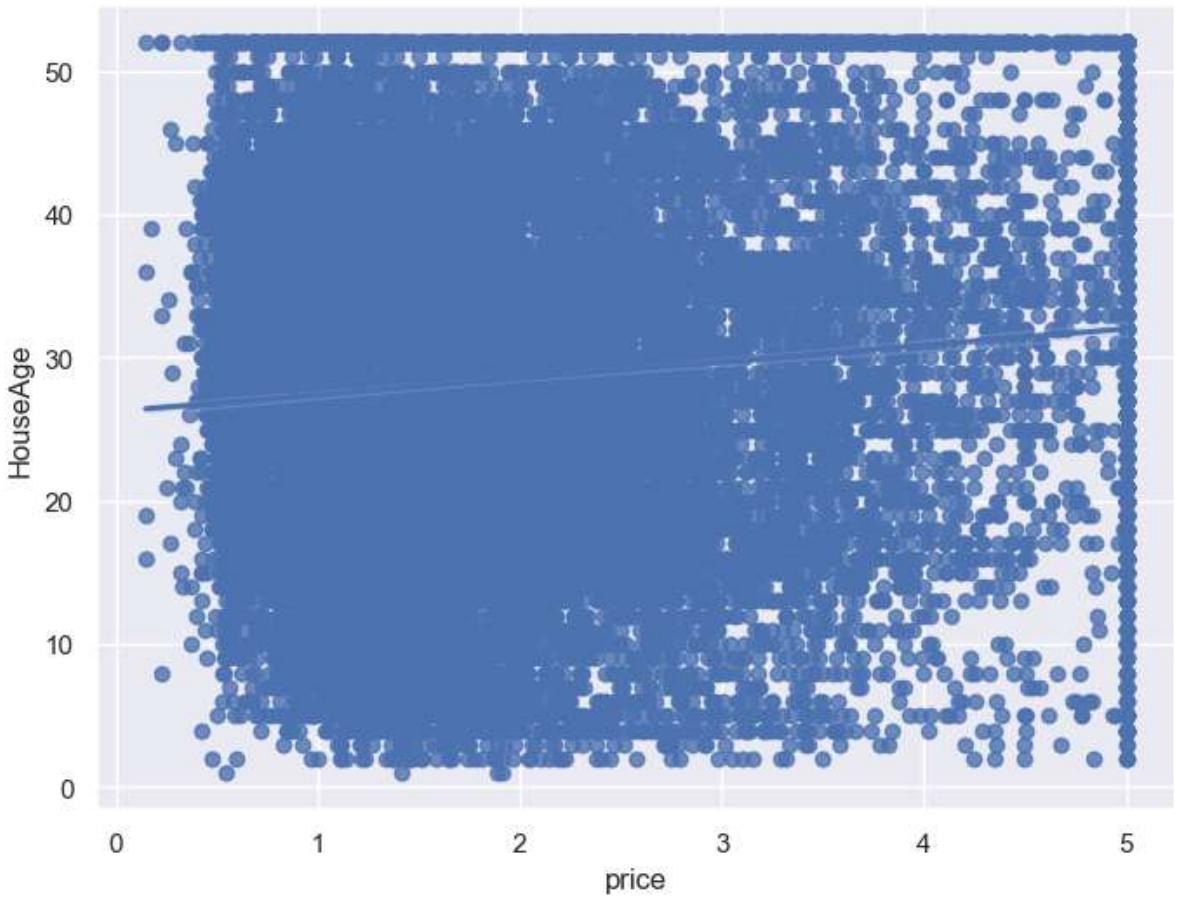
```
In [16]: sns.set(rc={'figure.figsize':(8,6)})  
sns.regplot(x='price',y= 'MedInc',data=dataset )
```

```
Out[16]: <AxesSubplot: xlabel='price', ylabel='MedInc'>
```



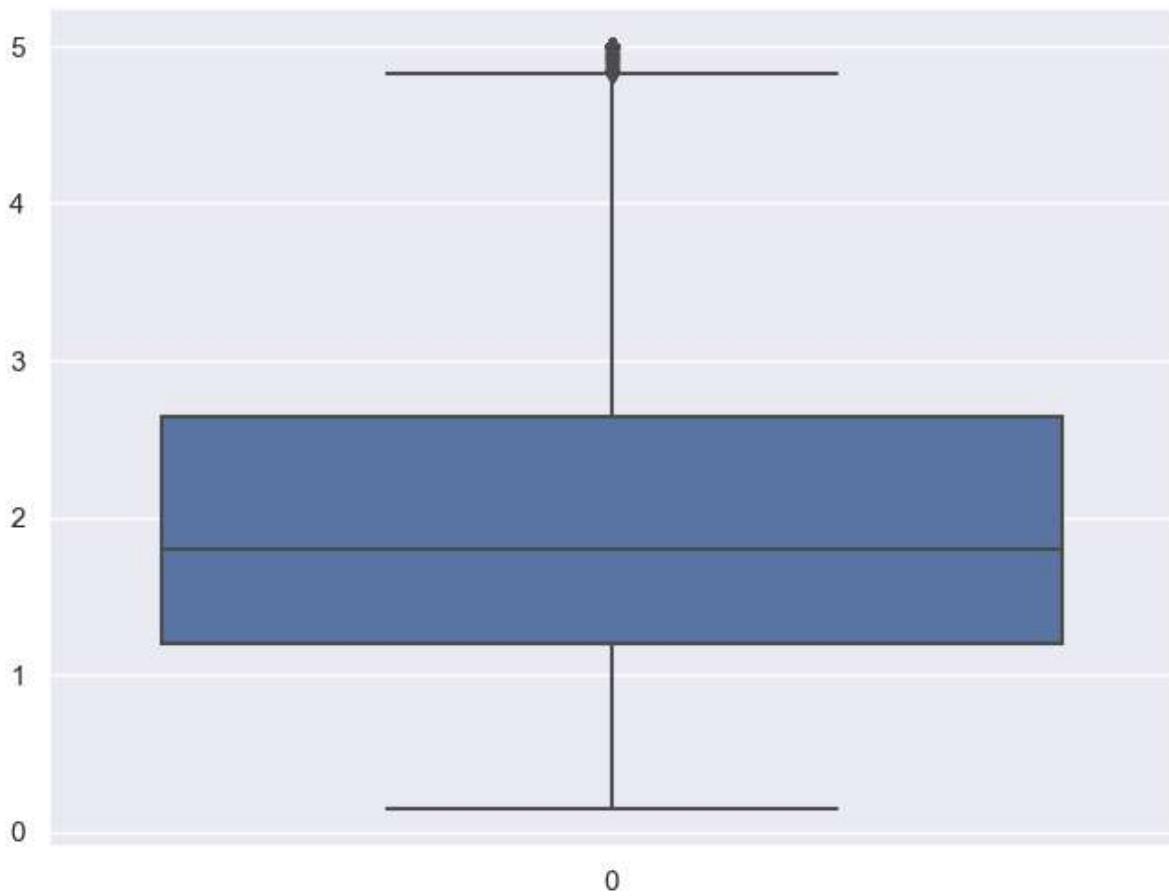
```
In [17]: sns.set(rc={'figure.figsize':(8,6)})  
sns.regplot(x='price',y= 'HouseAge',data=dataset )
```

```
Out[17]: <AxesSubplot: xlabel='price', ylabel='HouseAge'>
```



```
In [18]: sns.boxplot(dataset['price'])
```

```
Out[18]: <AxesSubplot: >
```



In [19]: *#independent and dependent feature*

```
x=dataset.iloc[:, :-1]  
y=dataset.iloc[:, -1]
```

In [20]: **x**

Out[20]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25
...
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	-121.09
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807	39.49	-121.21
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635	39.43	-121.22
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	-121.32
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981	39.37	-121.24

20640 rows × 8 columns

In [21]: y

Out[21]:

```
0      4.526
1      3.585
2      3.521
3      3.413
4      3.422
...
20635    0.781
20636    0.771
20637    0.923
20638    0.847
20639    0.894
Name: price, Length: 20640, dtype: float64
```

In [22]:

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=12)
```

In [23]:

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[23]:

```
((13828, 8), (6812, 8), (13828,), (6812,))
```

In [24]:

```
6812/(13828+6812)*100 #33% data is in testing and rest is in training
```

Out[24]:

```
33.00387596899225
```

In [25]:

```
from sklearn.preprocessing import StandardScaler
scalar=StandardScaler()
```

```
In [26]: scalar
```

```
Out[26]: ▾ StandardScaler
```

```
    StandardScaler()
```

```
In [27]: x_train=scalar.fit_transform(x_train)
```

```
In [28]: x_test=scalar.fit_transform(x_test)
```

```
In [29]: x_train
```

```
Out[29]: array([[-1.24904603,  0.42633018, -0.61320277, ..., -0.04338037,
       -0.75605028,  0.611307  ],
      [-0.14406409,  0.50577515,  0.18856354, ..., -0.05536191,
       -0.11982222,  0.31137802],
      [ 0.9692072 ,  0.3468852 ,  0.63641081, ...,  0.00476754,
       -1.39227835,  1.25615429],
      ...,
      [-0.93899219, -1.24201426, -0.09375134, ..., -0.07052394,
       1.65787385, -1.02830472],
      [-0.67617706, -0.76534442, -0.67423435, ..., -0.05918003,
       -0.69523436,  0.66629397],
      [ 0.25929815,  0.74411007,  0.26440744, ..., -0.05489583,
       -1.34081873,  1.25115547]])
```

```
In [30]: x_test
```

```
Out[30]: array([[ 0.24916919,  1.06101413, -0.08969363, ...,  0.26576233,
       -0.74243808,  0.71319515],
      [-0.54873699, -0.13125719, -0.62158373, ..., -0.56352247,
       0.95047089, -1.25738263],
      [ 0.36609347, -1.00558949,  0.04177291, ...,  0.01951513,
       -0.93470751,  0.92219582],
      ...,
      [-0.13653111,  0.50462084,  0.06073846, ...,  0.02910616,
       0.9176444 , -0.59554717],
      [-0.69692031, -1.64146753,  0.31454275, ..., -0.88159557,
       -0.89250203,  1.17598235],
      [ 0.20612989,  1.85586167,  0.00917042, ..., -0.75619168,
       1.05832936, -1.34695435]])
```

```
In [49]: from sklearn.linear_model import LinearRegression
```

```
In [50]: regression=LinearRegression()
```

```
In [51]: regression.fit(x_train,y_train)
```

```
Out[51]: ▾ LinearRegression
```

```
    LinearRegression()
```

```
In [52]: regression.coef_
```

```
Out[52]: array([ 0.8160491 ,  0.11835219, -0.25538202,  0.3034737 , -0.00345005,
   -0.04064189, -0.9039433 , -0.87297273])
```

```
In [53]: regression.intercept_
```

```
Out[53]: 2.06527526974255
```

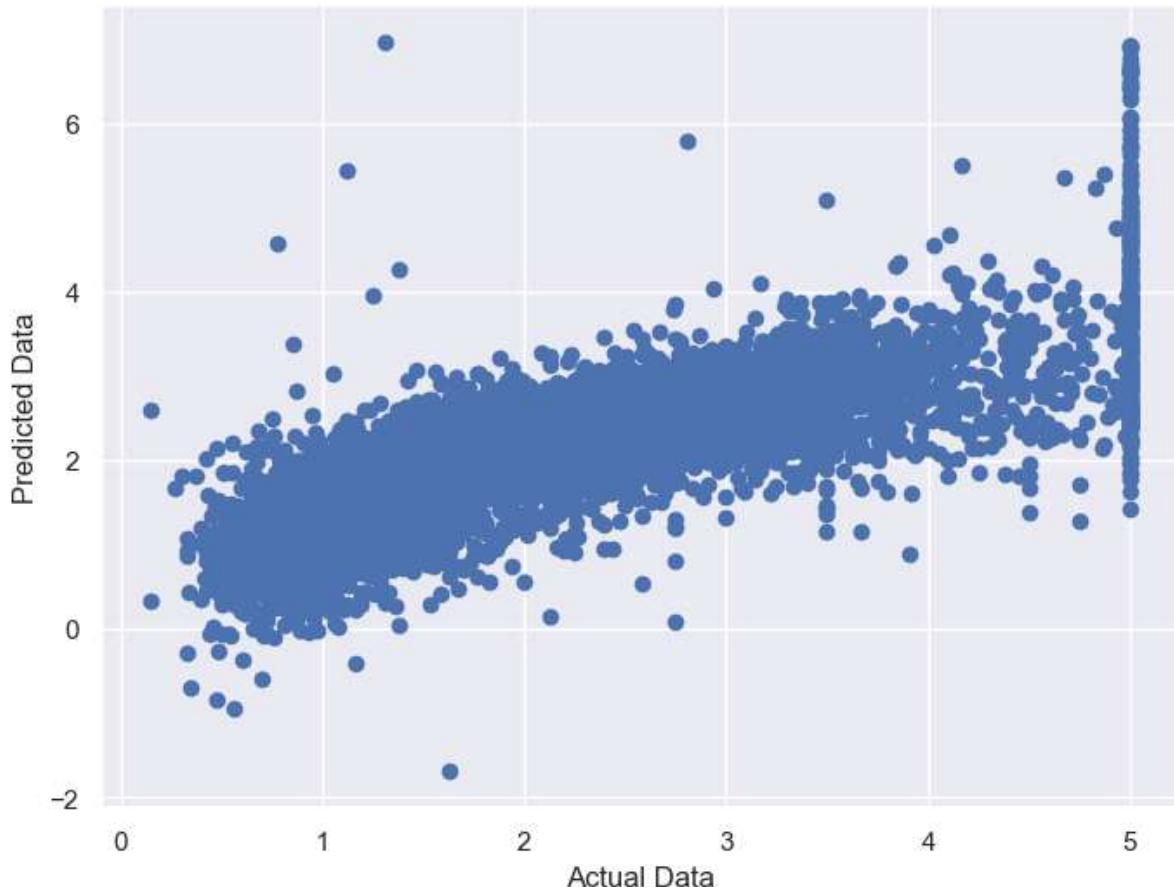
```
In [55]: reg_pred=regression.predict(x_test)
```

```
In [57]: print(reg_pred)
```

```
[2.43556627 2.06941919 2.28208702 ... 1.64898575 1.11214087 2.66106521]
```

```
In [65]: plt.scatter(y_test,reg_pred)
plt.xlabel('Actual Data')
plt.ylabel('Predicted Data')
```

```
Out[65]: Text(0, 0.5, 'Predicted Data')
```



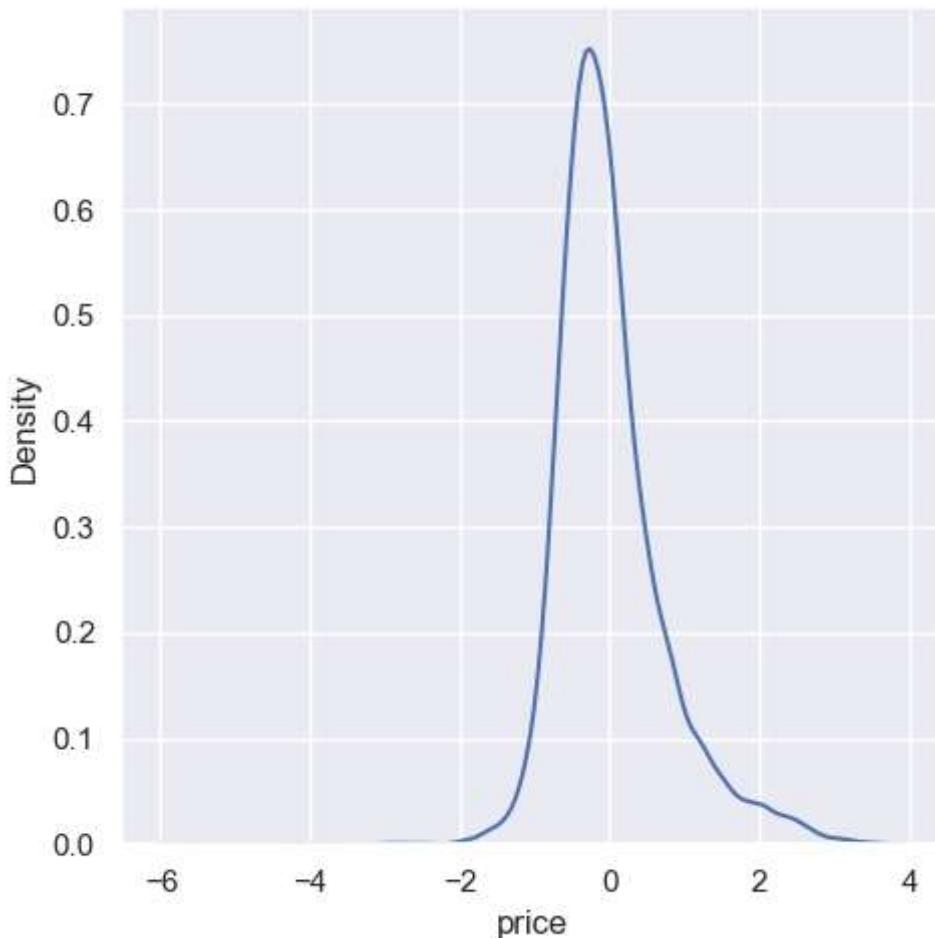
```
In [67]: residual=y_test-reg_pred
```

```
In [68]: residual
```

```
Out[68]: 6906    -0.321566
  767    -0.117419
 10555    0.135913
 17456   -1.391623
 20617   -0.379597
      ...
  861    -0.096177
 10326   -0.022179
 19588    0.137014
 12146   -0.011141
  425     0.187935
Name: price, Length: 6812, dtype: float64
```

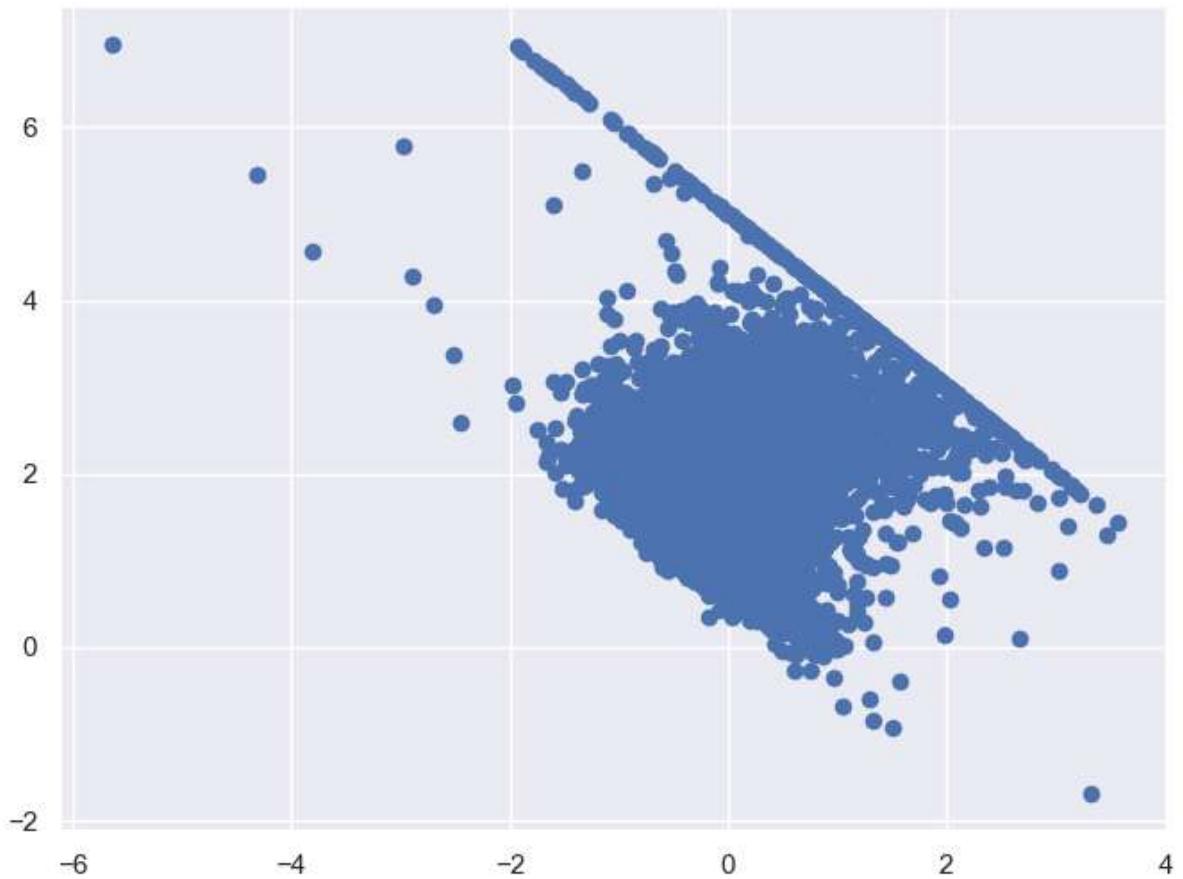
```
In [69]: sns.displot(residual, kind='kde')
```

```
Out[69]: <seaborn.axisgrid.FacetGrid at 0x244f44bf010>
```



```
In [71]: plt.scatter(residual, reg_pred)
```

```
Out[71]: <matplotlib.collections.PathCollection at 0x244f4779f00>
```



Performance Metrics

```
In [77]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test, reg_pred))
print(mean_absolute_error(y_test, reg_pred))
print(np.sqrt(mean_squared_error(y_test, reg_pred)))
```

0.5204783529558654
0.5257562312164743
0.7214418569475058

R Square Value and Adjusted R Square

```
In [78]: from sklearn.metrics import r2_score
score=r2_score(y_test,reg_pred)
```

```
In [79]: score
```

```
Out[79]: 0.6125609410014536
```

```
In [80]: # Adjusted R Square
1-(1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[80]: 0.6121053313480671
```

```
In [127]: df1=pd.DataFrame({'Actual':y_test,'Predict':reg_pred})  
df1
```

```
Out[127]:
```

	Actual	Predict
6906	2.114	2.435566
767	1.952	2.069419
10555	2.418	2.282087
17456	1.283	2.674623
20617	0.708	1.087597
...
861	2.169	2.265177
10326	2.832	2.854179
19588	1.786	1.648986
12146	1.101	1.112141
425	2.849	2.661065

6812 rows × 2 columns

Lasso Regression Model

```
In [82]: from sklearn.linear_model import Lasso  
lasso = Lasso()  
lasso
```

```
Out[82]:
```

▼ Lasso
Lasso()

```
In [84]: lasso.fit(x_train, y_train)
```

```
Out[84]:
```

▼ Lasso
Lasso()

```
In [86]: print(lasso.coef_)
```

[0. 0. 0. -0. -0. -0. -0. -0.]

```
In [92]: print(lasso.intercept_)
```

2.065275269742551

```
In [93]: lasso_predict=lasso.predict(x_test)
```

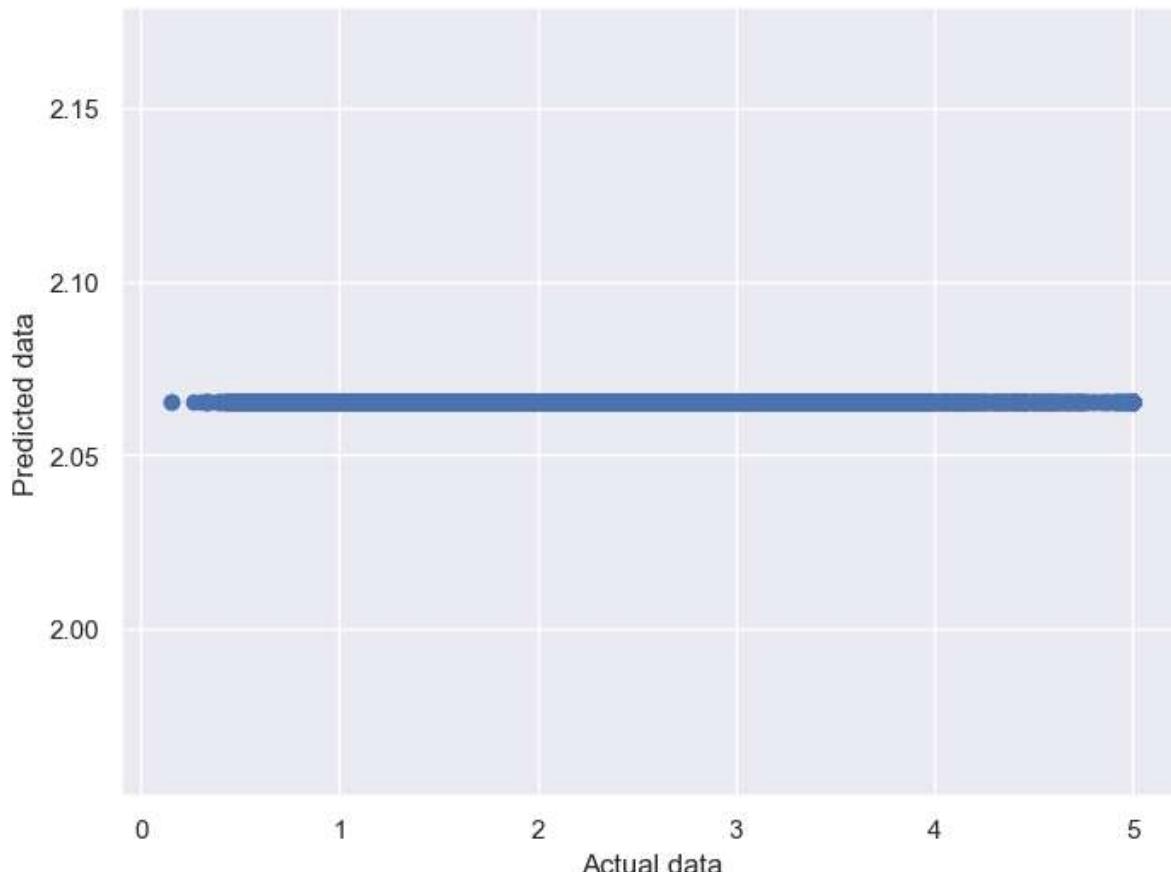
```
In [94]: lasso_predict
```

```
Out[94]: array([2.06527527, 2.06527527, 2.06527527, ..., 2.06527527, 2.06527527, 2.06527527])
```

Assumption of lasso

```
In [97]: plt.scatter(y_test,lasso_predict)  
plt.xlabel('Actual data')  
plt.ylabel('Predicted data')
```

```
Out[97]: Text(0, 0.5, 'Predicted data')
```



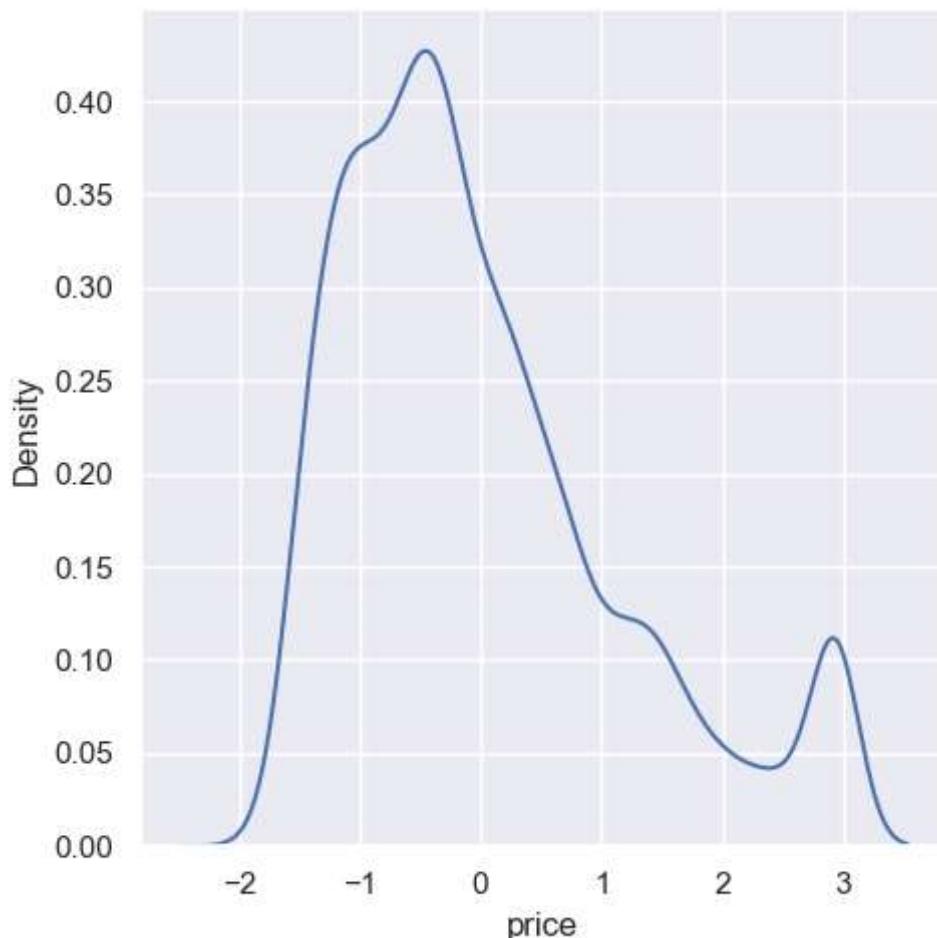
```
In [98]: residual = y_test-lasso_predict
```

```
In [99]: residual
```

```
Out[99]: 6906    0.048725
767     -0.113275
10555   0.352725
17456   -0.782275
20617   -1.357275
...
861     0.103725
10326   0.766725
19588   -0.279275
12146   -0.964275
425     0.783725
Name: price, Length: 6812, dtype: float64
```

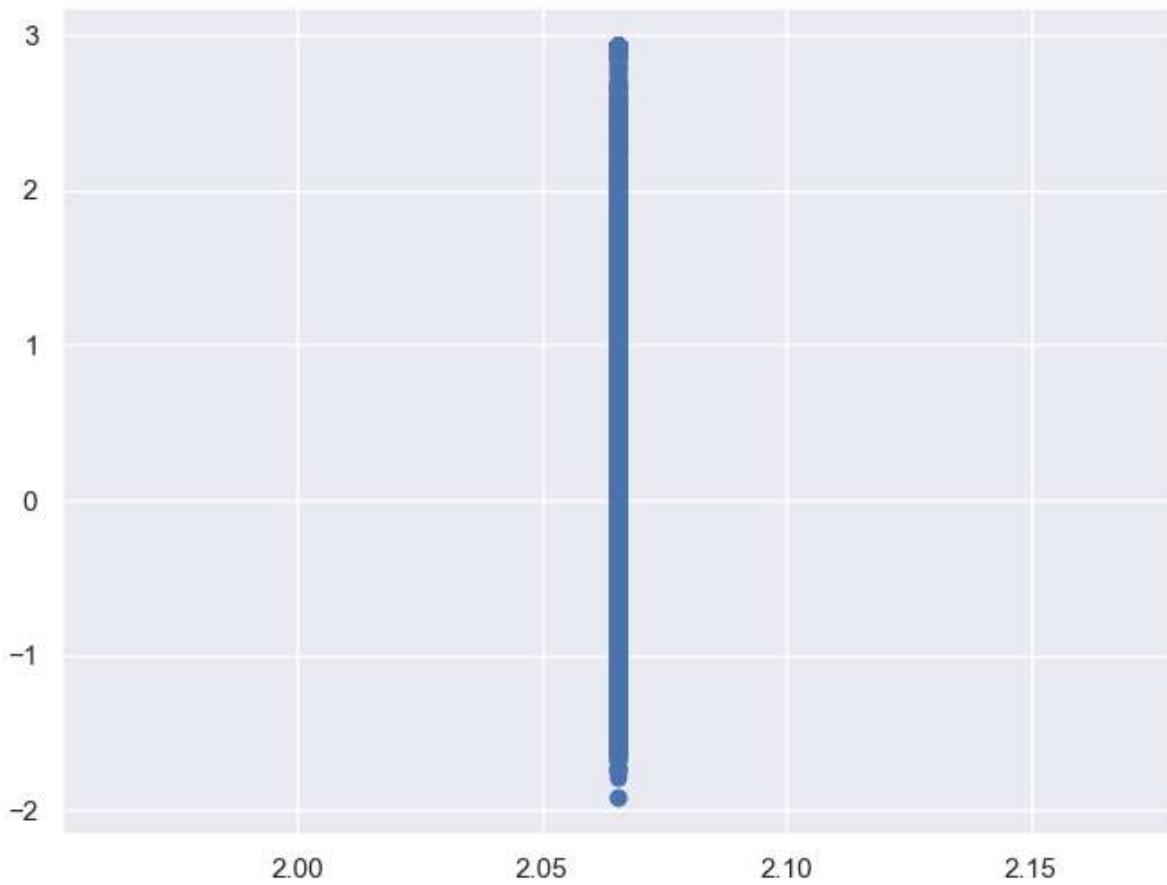
```
In [100... sns.displot(residual,kind='kde')
```

```
Out[100]: <seaborn.axisgrid.FacetGrid at 0x244f40e8520>
```



```
In [101... plt.scatter(lasso_predict,residual)
```

```
Out[101]: <matplotlib.collections.PathCollection at 0x244f4dbc550>
```



Performance Metrics

In [102...]

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,lasso_predict))
print(mean_absolute_error(y_test,lasso_predict))
print(np.sqrt(mean_squared_error(y_test,lasso_predict)))
```

1.3434801558184395

0.9136594481730159

1.1590859139073513

R Square Value and Adjusted R Square

In [103...]

```
from sklearn.metrics import r2_score
score=r2_score(y_test,lasso_predict)
score
```

Out[103]: -7.365224208366605e-05

In [104...]

```
# Adjusted R Square
1-(1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[104]: -0.0012496906395460528
```

```
In [126... df1=pd.DataFrame({'Actual':y_test,'Predict':lasso_predict})  
df1
```

```
Out[126]:
```

	Actual	Predict
6906	2.114	2.065275
767	1.952	2.065275
10555	2.418	2.065275
17456	1.283	2.065275
20617	0.708	2.065275
...
861	2.169	2.065275
10326	2.832	2.065275
19588	1.786	2.065275
12146	1.101	2.065275
425	2.849	2.065275

6812 rows × 2 columns

Elasticnet Regression Model

```
In [106... from sklearn.linear_model import ElasticNet  
  
elastic=ElasticNet()  
elastic
```

```
Out[106]:
```

```
▼ ElasticNet  
ElasticNet()
```

```
In [107... elastic.fit(x_train,y_train)
```

```
Out[107]:
```

```
▼ ElasticNet  
ElasticNet()
```

```
In [108... print(elastic.coef_)
```

```
[ 0.19553816  0.          0.          -0.          -0.          -0.  
 -0.          -0.        ]
```

```
In [109... print(elastic.intercept_)
```

```
2.065275269742551
```

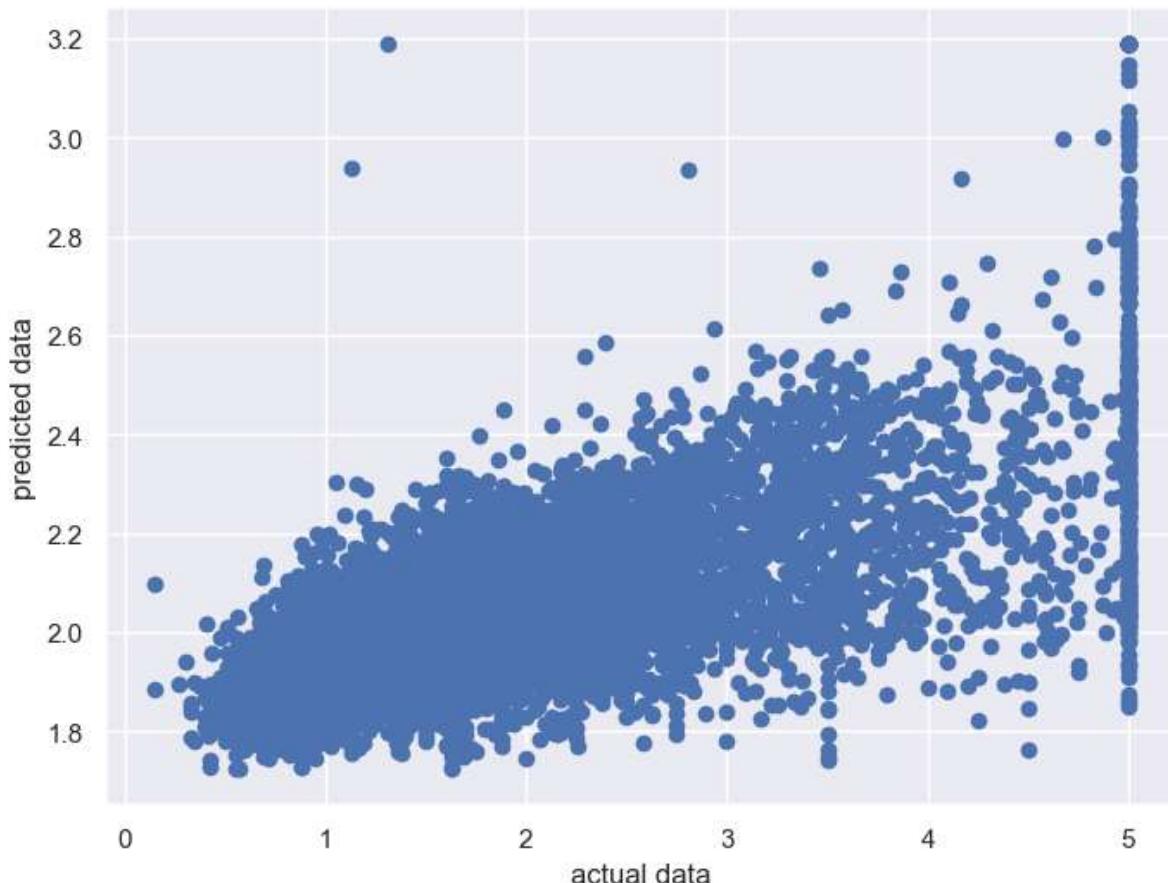
```
In [112]: elastic_pred=elastic.predict(x_test)  
elastic_pred
```

```
Out[112]: array([2.11399735, 1.95797625, 2.13686051, ..., 2.03857823, 1.92900076,  
2.10558153])
```

Assumption of Elasticnet Regression model

```
In [116]: plt.scatter(y_test,elastic_pred)  
plt.xlabel('actual data')  
plt.ylabel('predicted data')
```

```
Out[116]: Text(0, 0.5, 'predicted data')
```



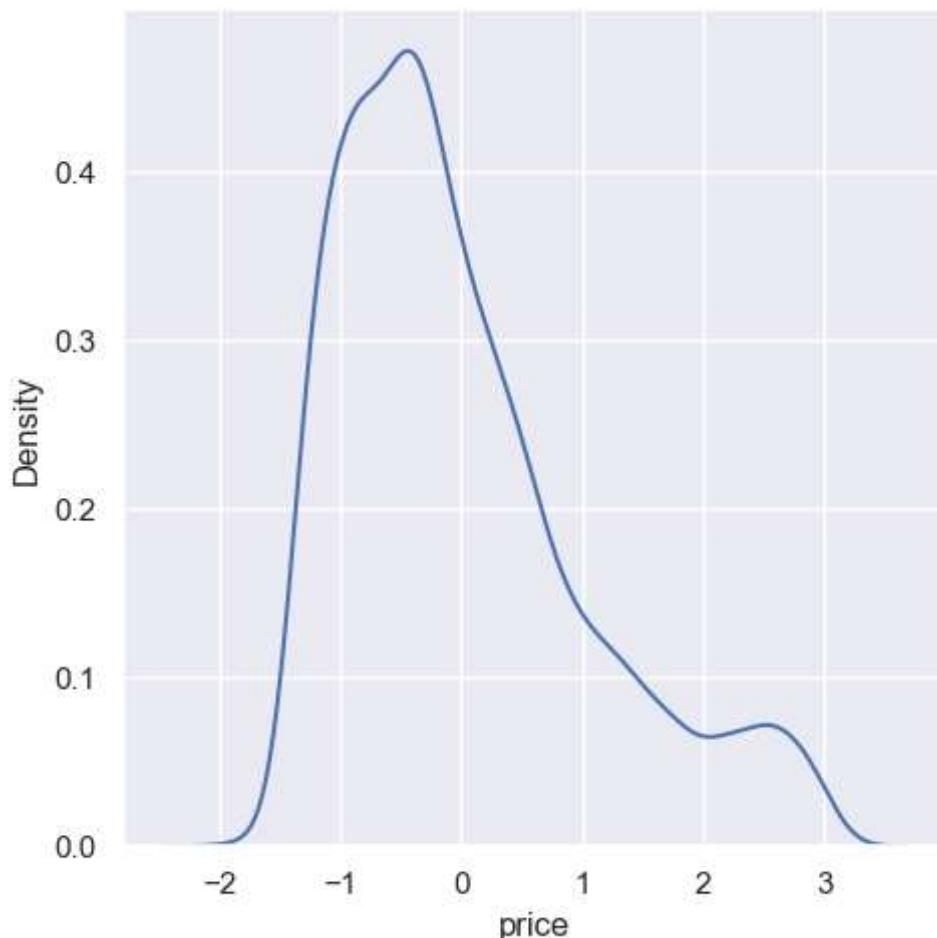
```
In [117]: residual = y_test-elastic_pred
```

```
In [118]: residual
```

```
Out[118]: 6906    0.000003
  767   -0.005976
 10555   0.281139
 17456   -0.753982
 20617   -1.298653
      ...
  861    0.123845
 10326   0.485087
 19588   -0.252578
 12146   -0.828001
  425    0.743418
Name: price, Length: 6812, dtype: float64
```

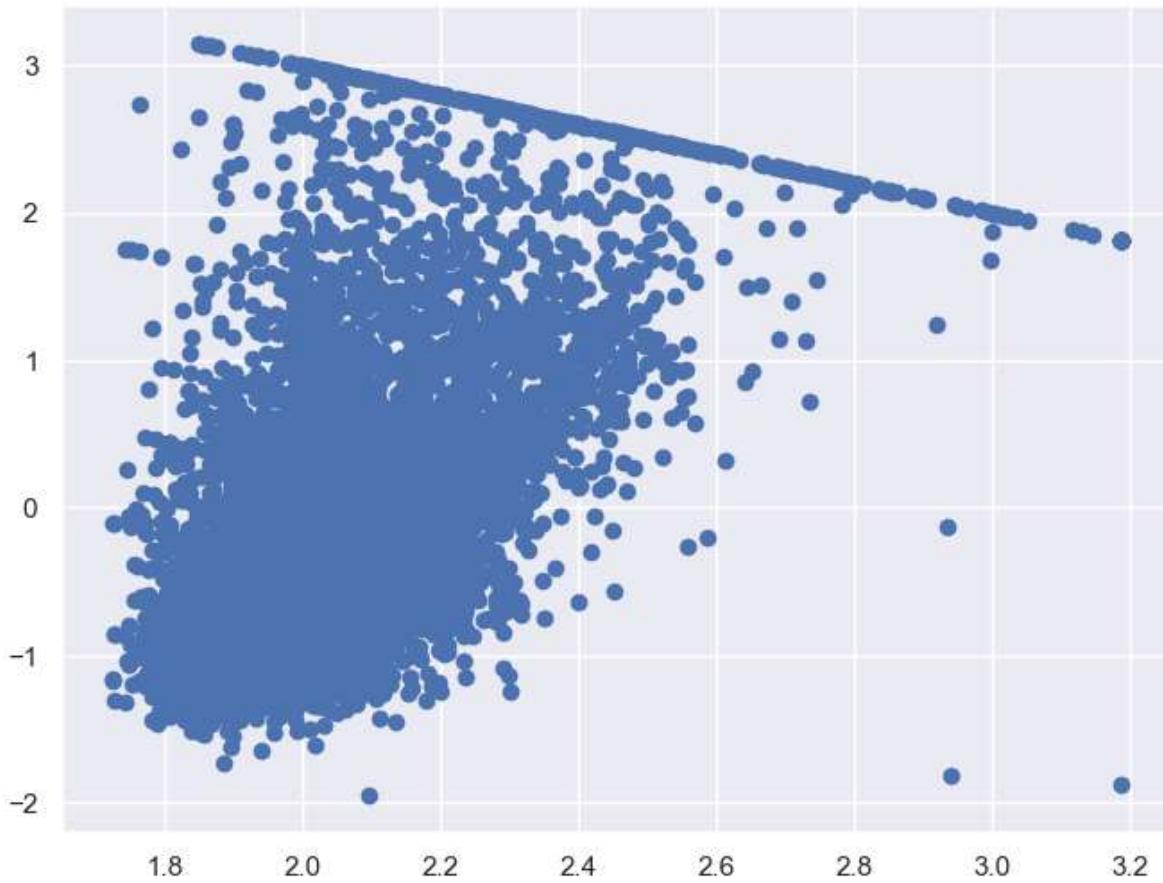
```
In [119... sns.displot(residual,kind='kde')
```

```
Out[119]: <seaborn.axisgrid.FacetGrid at 0x244f4e4c250>
```



```
In [120... plt.scatter(elastic_pred,residual )
```

```
Out[120]: <matplotlib.collections.PathCollection at 0x244f4dbcb50>
```



Perfomance Metrics

In [121]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,elastic_pred))
print(mean_absolute_error(y_test,elastic_pred))
print(np.sqrt(mean_squared_error(y_test,elastic_pred)))
```

1.0706121914761009

0.8144919283598954

1.0347039148839154

R Square and Adjusted R Square

In [123]:

```
from sklearn.metrics import r2_score
score=r2_score(y_test,elastic_pred)
score
```

Out[123]: 0.2030466249706926

In [124]:

```
# Adjusted R Square
1-(1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[124]: 0.2021094462259866
```

```
In [125... df1=pd.DataFrame({'Actual':y_test,'Predict':elastic_pred})  
df1
```

```
Out[125]:
```

	Actual	Predict
6906	2.114	2.113997
767	1.952	1.957976
10555	2.418	2.136861
17456	1.283	2.036982
20617	0.708	2.006653
...
861	2.169	2.045155
10326	2.832	2.346913
19588	1.786	2.038578
12146	1.101	1.929001
425	2.849	2.105582

6812 rows × 2 columns

```
In [ ]:
```